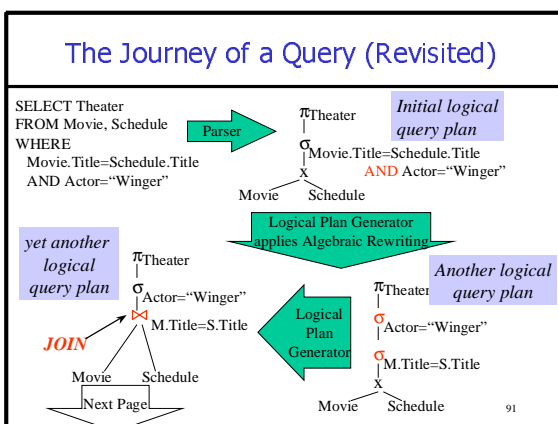


Query Processor

- Basic task: turn user queries (and updates) into a sequence of operations (a *plan*) on the database
=> from high level queries to low level commands
- Many choices:
 - Which of the **algebraically equivalent** forms of a query will lead to the most efficient algorithm?
 - For each algebraic operator which **algorithm** should we use to run the operator?
 - How should the operations pass data from one to the other? (eg. main memory buffers, disk buffers)

90



The Journey of a Query (Revisited)

$\pi_{Theater}$
 $\sigma_{Actor="Winger"}$
 \Join
 Movie Schedule

Logical Plan Generator

4th logical query plan

$\pi_{Theater}$
 \Join
 $\sigma_{Actor="Winger"}$
 Schedule Movie

σ_{cond}
 \Join
 σ_{cond}
 R S

if cond refers only to S

- 4 logical query plans created
- algebraic rewritings* were used to produce candidate logical query plans
- the last one wins
- multiple logical plans may “win” eventually (and be considered subsequently)

92

The Journey Continues at the Physical Plan Generator

$\pi_{Theater}$
 \Join
 $\sigma_{Actor="Winger"}$
 Schedule Movie

Physical Plan Generator

Physical Plan Generator chooses execution primitives and data passing

index on Actor, tables Schedule sorted on Title,

$\pi_{Theater}$
 \Join
 $\sigma_{Actor="Winger"}$
 Schedule Movie

Physical Plan 2

index on Actor and Title, unsorted tables, tables >> memory

Physical Plan 1

More than one plan may be generated by choosing different primitives 93

More Than One Plan May be Generated and Considered for Evaluation

\Join
 $\sigma_{R.A=T.A}$
 R T

Transform

\Join
 $\sigma_{R.A=S.A}$
 R S

Transform

\Join
 $\sigma_{R.A=T.A}$
 R T

94

Issues in Query Processing and Optimization

- *Generate Plans*
 - systematically transform expressions
 - employ execution primitives for computing relational algebra operations
- *Estimate Cost of Generated Plans*
 - Statistics
- “Smart” Search of the Space of Possible Plans
 - always do the “good” transformations (relational algebra optimization)
 - prune the space (e.g., System R)
- Often the above steps are mixed

95

Bag Semantics of Operators

- SQL semantics for collections of tuples:
 - often **bags** (aka **multisets**) instead of sets.
 - **Union(R,S)**: a tuple t occurs in the result as many times as the **sum** of occurrences in each of R and S
 - **Intersection of R and S** : ... the **minimum** number of occurrences in R and S
 - **Difference of R and S** : ... occurrences in R **minus** occurrences in S
 - $\delta(R)$ converts the **bag** R into a **set**
 - SQL’s **R UNION S** by default $\delta(R \cup S)$
- **Example:** $R = \{A, B, B\}$ and $S = \{C, A, B, C\}$.

96

Extended Algebra Operators

- **selection** $\sigma_C(R)$: condition C over arithmetic-, string-, and comparison-operators
 - WHERE in SQL, but C depends only on individual tuples
- **projection** $\pi_A(R)$: attribute list A may include
 - $x \rightarrow y$ (rename x to y in the output schema)
 - $a+b \rightarrow x$ (output schema contains x column with sum of a and b)
 - $c||d \rightarrow y$ (concatenate c and d as y)
 - again: look at one tuple at a time
- alternatively: **special purpose operators**:
 - **MULT** _{$A, B \rightarrow C$} (R):
 - for each tuple of R , multiply attribute A with attribute B and put the result in a new attribute named C .
 - **PLUS** _{$A, B \rightarrow C$} (R), **CONCAT** _{$A, B \rightarrow C$} (R), ...

97

Product and Joins

- **product ($R \times S$):**
 - r is n times in R , s is m times in S , then (r,s) is $n*m$ times in the product
 - (qualify attributes if necessary: $R.a$ vs $S.a$)
- **natural join $R \bowtie S = \pi_A \sigma_C(R \times S)$ where**
 - C equates all common attributes
 - A has all attributes from R and S (but only one copy per equated attribute)
 - (definition of join ~ simple rewriting rule)
- **theta join**
 - arbitrary condition involving attributes

98

Grouping and Aggregation


- **Grouping & Aggregation in SQL:**

```
SELECT selection-list FROM from-list WHERE...
GROUP BY groupby-list
HAVING group-qualification
```
- **Example:** given **movie(title, year, actor)**, find for each actor having more than two movies, the first year of appearance

```
SELECT actor, MIN(year) as firstYear FROM movie
GROUP BY actor
HAVING COUNT(title) > 2
```
- Every attribute in the *selection-list* must appear in the *groupby-list* (*Why?*)
- $grP_{actor, MIN(year) \Rightarrow firstYear, COUNT(title) \Rightarrow aux}(movie)$

99

Algebraic Rewritings: Commutativity and Associativity



- Do the above hold for both sets and bags?
- Do commutativity and associativity hold for arbitrary Theta Joins?

100

Algebraic Rewritings: Commutativity and Associativity

